

ADARSH KUMAR SINGH

AI/ML Engineer | MLOps | Data Scientist

✉ ultraman5115@gmail.com ☎ +91-9439440544 📍 Greater Delhi Area, India 🔗 ai-apex-paragon-qcls-beta.vercel.app

🌐 linkedin.com/in/i-am-entity 🏠 github.com/Cyber-Duelist 🏠 kaggle.com/ultramanfindeye 🏠 huggingface.co/spaces/IamEntity

🇮🇳 Indian

📄 SUMMARY

AI/ML Engineer with hands-on experience building production-grade multi-agent systems, RAG pipelines, and MLOps dashboards. Designed autonomous DevOps swarms using LLaMA 3 that reduced manual debugging by 90%. Oracle Certified in Generative AI and Data Science. Skilled in Python, FastAPI, LangChain, ChromaDB, Docker, and CI/CD automation. Seeking AI/ML Engineer, MLOps Engineer, or Data Scientist roles.

📁 PROFESSIONAL EXPERIENCE

AI Software Engineer, Self-Employed

01/2026 – Present | Remote

AI-Apex-Paragon — Autonomous Systems Lab

- Designed and shipped 8 production-grade AI systems in a 14-week intensive engineering sprint
- Deployed 3 live demos on HuggingFace Spaces and 2 on Vercel with public-facing endpoints
- Authored technical articles on multi-agent architecture and RAG pipeline optimization

AI Systems Member,

01/2024 – 05/2026

Software Incubator Lab — Dr. A.P.J. Abdul Kalam Technical University

- Built a 4-agent DevOps swarm reducing mean-time-to-recovery by 90% on CI/CD failures
- Deployed containerized AI microservices using Docker and Kubernetes on AWS and OCI
- Configured CI/CD pipelines with GitHub Actions for automated testing and deployment

Software Development Lead,

08/2023 – 05/2026

Programming Club — Dr. A.P.J. Abdul Kalam Technical University

- Architected a RAG pipeline ingesting enterprise PDFs with 94% retrieval accuracy
- Reduced LLM hallucination by 78% using semantic chunking and reranking guardrails
- Mentored 15+ juniors on Python best practices, FastAPI architecture, and Git workflows

🎓 EDUCATION

Bachelor of Technology — Electrical, Electronics & Communications Engineering, Dr. A.P.J. Abdul Kalam Technical University

2022 – 2026

🧠 TECHNICAL SKILLS

Languages

Python, SQL, JavaScript, Bash

AI/ML

LLaMA 3, GPT, RAG, LangChain, Prompt Engineering, Multi-Agent Orchestration, ReAct Loops

Data Science

Pandas, NumPy, scikit-learn, XGBoost, Random Forest, Gradient Boosting, Matplotlib, Seaborn, EDA, Feature Engineering, Hyperparameter Tuning, Cross-Validation

MLOps

MLflow, Model Failover Routing, LLM-as-a-Judge Evaluation, Guardrails (Input/Output), Structured Logging

Databases

ChromaDB, FAISS, PostgreSQL, Microsoft SQL Server, Vector Databases

Frameworks

FastAPI, Streamlit, Next.js, Sentence-Transformers

Cloud & DevOps

Docker, Kubernetes, AWS (EC2, S3), Oracle Cloud (OCI), GitHub Actions, CI/CD, Vercel

Tools

Git, REST APIs, WebSockets, JSON Schema, Pytest

📜 CERTIFICATIONS

Oracle Generative AI Professional

Issuer: Oracle Date: Oct 2025 Credential: 3229858040C125GAIOCP

Oracle Data Science Professional

Issuer: Oracle Date: Oct 2025 Credential: 3229858040C125DSOCP

Artificial Intelligence — 97th Percentile

Issuer: TestGorilla Date: Jun 2026 Credential: Verified

Microsoft SQL Server — 99th Percentile

Issuer: TestGorilla Date: Jun 2026 Credential: Verified

Kaggle — Data Science & Machine Learning

Issuer: Kaggle Date: 2025 Credential: Verified

PROJECTS

CLARA AI — Clinical Language & Adaptive Routing Assistant, *Next.js · Groq LLaMA 3.3 · Vision AI · WebSockets · TTS*

- Built a multimodal AI receptionist automating 85% of triage routing with 99.9% uptime during load testing
- Integrated Vision AI for real-time wound analysis and auto-updating Clinical Command Center
- Implemented Groq Whisper STT and TTS for bidirectional voice interaction with sub-500ms TTFB

Autonomous Self-Healing DevOps Swarm, *Python · Groq · LLaMA 3 · Pytest · Multi-Agent Orchestration*

- Built a 4-agent swarm intercepting CI/CD failures, diagnosing logs, and auto-generating patches
- Reduced mean-time-to-recovery by 90% using ReAct reasoning loops and sandboxed testing
- Achieved 85% first-pass fix accuracy across 50+ automated test scenarios

Enterprise Production Agent — LLMOps Dashboard, *Python · FastAPI · Guardrails · LLMOps*

- Developed a hardened compliance agent with Input/Output Guardrails blocking 100% of prompt injections
- Built an LLM-as-a-Judge Evaluation Dashboard comparing 4 model architectures in real time
- Implemented Dynamic Model Failover routing ensuring 99.9% uptime during API rate-limit spikes

PersonaDoc — Document Intelligence RAG, *RAG · ChromaDB · Embeddings · FastAPI*

- Architected a RAG pipeline ingesting PDFs with 94% retrieval accuracy using Sentence-Transformers
- Reduced LLM hallucination by 78% via semantic chunking, reranking, and citation guardrails
- Exposed 3 FastAPI REST endpoints processing 96% extraction accuracy in 8 seconds per document

PUBLICATIONS

Building an Autonomous Self-Healing CI/CD Swarm

Production RAG: Beating Hallucinations with Semantic Reranking