

ADARSH KUMAR SINGH

Ghaziabad, Uttar Pradesh | +91-94394-40544 | adarshentity098@gmail.com

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Results-driven Full Stack AI Software Engineer specializing in Large Language Models (LLMs), Agentic Workflows, and scalable Python architectures. Proven ability to engineer autonomous multi-agent systems, implement Retrieval-Augmented Generation (RAG) pipelines, and build robust enterprise LLM guardrails from the ground up. Seeking an AI Engineering role to leverage expertise in Python, FastAPI, and GenAI to build secure, highly automated applications.

SKILLS

Programming Languages: Python, JavaScript, SQL, Bash

AI & Machine Learning: Large Language Models (LLaMA 3, GPT), Retrieval-Augmented Generation (RAG), Multi-Agent Orchestration (ReAct Loops), AI Guardrails, Prompt Engineering

Frameworks & Libraries: FastAPI, Pytest, Sentence-Transformers, ChromaDB (Vector Databases), GSAP, Three.js

Cloud & Infrastructure: Oracle Cloud Infrastructure (OCI), Oracle Database Design, Data Warehousing, Server-Sent Events (SSE)

Developer Tools: Git, REST APIs, JSON Validation, CI/CD Pipelines

PROJECTS

Autonomous Self-Healing DevOps Swarm

Multi-Agent Orchestration | Python, LLaMA 3, Pytest, GitHub API | Mar 2026 - Present

- Engineered a multi-agent orchestration framework from scratch using ReAct (Reason + Act) loops to intercept CI/CD pipeline failures and automatically remediate broken codebase branches, reducing manual debugging time by 90%.
- Implemented a continuous diagnostic loop utilizing LLaMA 3 to analyze raw error logs, inject Python patches via custom file operations, and verify fixes using local sandboxed Pytest environments.
- Designed 4 dynamic agent personas (Diagnoser, Developer, Verifier, PR Reviewer) achieving zero-touch bug remediation across 100% of tested workflows with seamless GitHub API integration.

Enterprise Compliance Agent + LLMops Dashboard

Agentic Workflow | Python, FastAPI, Groq, GitHub Actions | Jan 2026 - Feb 2026

- Developed a production-grade compliance AI agent equipped with persistent state memory, tool calling, and rigorous LLM security boundaries for enterprise deployment.
- Designed pre-execution Input Guardrails to block adversarial prompt injection and post-execution Output Guardrails to enforce strict domain-scope adherence on all LLM responses.
- Built an automated LLM-as-a-Judge Evaluation Dashboard with Dynamic Model Failover routing (70B to 8B), guaranteeing 99.9% uptime during API rate-limit bottlenecks.

PersonaDoc - Production RAG System

Retrieval-Augmented Generation | FastAPI, ChromaDB, Sentence-Transformers, LLaMA 3.3 | Nov 2025 - Dec 2025

- Architected a Retrieval-Augmented Generation (RAG) service capable of ingesting enterprise PDFs and generating cited, hallucination-free answers using semantic search and reranking.
- Exposed the service via a FastAPI REST backend with /upload, /search, and /documents endpoints, utilizing semantic chunking strategies and ChromaDB vector database for high-context retrieval.

AI Code Reviewer + Meeting Intelligence System

LLM Microservices | Python, FastAPI, GitHub REST API, Groq | Sep 2025 - Oct 2025

- Built a microservice that connects to the GitHub API, fetches raw pull request diffs, and streams changes to a high-reasoning LLM for deep architectural code review with structured JSON output.
- Developed a two-stage Meeting Intelligence pipeline that ingests raw transcripts and outputs structured summaries, per-person action items with deadlines, and ready-to-send follow-up emails.

EDUCATION

Dr. Abdul Kalam Technical University

B. Tech in Electronics and Communication Engineering | Aug 2022 - May 2026 | CGPA: 7.9

CERTIFICATIONS

- Oracle GenAI Certified Professional (2025)
- Oracle Data Science Professional Certification (2025)
- IoT and Industrial Automation Certification (2025)